

Homework 4

Kate Miller

Table of contents

.....	4
Question 1	4
Question 2	11
Question 3	14
70 points	14
and making sure that the accuracy is 100% while the errors are 0%.	15

! Important

Please read the instructions carefully before submitting your assignment.

1. This assignment requires you to only upload a **PDF** file on Canvas
2. Don't collapse any code cells before submitting.
3. Remember to make sure all your code output is rendered properly before uploading your submission.

Please add your name to the author information in the frontmatter before submitting your assignment

We will be using the following libraries:

```
packages <- c(
  "dplyr",
  "readr",
  "tidyr",
  "purrr",
  "stringr",
```

```
"corrplot",  
"car",  
"caret",  
"torch",  
"nnet",  
"broom"  
)  
  
#renv::install(packages)  
  
sapply(packages, require, character.only=T)
```

Loading required package: dplyr

Attaching package: 'dplyr'

The following objects are masked from 'package:stats':

filter, lag

The following objects are masked from 'package:base':

intersect, setdiff, setequal, union

Loading required package: readr

Loading required package: tidyr

Loading required package: purrr

Loading required package: stringr

Loading required package: corrplot

corrplot 0.92 loaded

Loading required package: car

Loading required package: carData

Attaching package: 'car'

The following object is masked from 'package:purrr':

some

The following object is masked from 'package:dplyr':

recode

Loading required package: caret

Loading required package: ggplot2

Loading required package: lattice

Attaching package: 'caret'

The following object is masked from 'package:purrr':

lift

Loading required package: torch

Warning: package 'torch' was built under R version 4.3.3

Loading required package: nnet

Warning: package 'nnet' was built under R version 4.3.3

Loading required package: broom

Warning: package 'broom' was built under R version 4.3.3

dplyr	readr	tidyr	purrr	stringr	corrplot	car	caret
TRUE	TRUE	TRUE	TRUE	TRUE	TRUE	TRUE	TRUE
torch	nnet	broom					
TRUE	TRUE	TRUE					

Question 1

💡 30 points

Automatic differentiation using `torch`

1.1 (5 points)

Consider $g(x, y)$ given by

$$g(x, y) = (x - 3)^2 + (y - 4)^2.$$

Using elementary calculus derive the expressions for

$$\frac{d}{dx}g(x, y), \quad \text{and} \quad \frac{d}{dy}g(x, y).$$

Answer: $d/dx \, g(x, y) = 2(x-3)$ $d/dy \, g(x, y) = 2(y-4)$ Partial derivative with respect to y :

Using your answer from above, what is the answer to

$$\left. \frac{d}{dx}g(x, y) \right|_{(x=3, y=4)} \quad \text{and} \quad \left. \frac{d}{dy}g(x, y) \right|_{(x=3, y=4)} \quad ?$$

Answer: $d/dx \, g(x, y) = 2(x-3)$ $d/dy \, g(x, y) = 2(y-4)$

Define $g(x, y)$ as a function in R, compute the gradient of $g(x, y)$ with respect to $x = 3$ and $y = 4$. Does the answer match what you expected?

```
g <- function(x, y) {  
  return((x - 3)^2 + (y - 4)^2)  
}  
gradient_x <- function(x, y) {  
  return(2 * (x - 3))  
}  
gradient_y <- function(x, y) {  
  return(2 * (y - 4))  
}  
x_value <- 3  
y_value <- 4  
  
gradient_x_at_3_4 <- gradient_x(x_value, y_value)  
gradient_y_at_3_4 <- gradient_y(x_value, y_value)
```

```
gradient_x_at_3_4
```

[1] 0

```
gradient_y_at_3_4
```

[1] 0

Yes, the answer matches what I was expecting.

1.2 (10 points)

Using elementary calculus derive the expressions for the gradients

Derived Expression: Please see images below.

calculation for $n=10$:

$$\frac{d}{du_{10}} h(u,v) = 3(u \cdot v)^2 v_{10} = 3(3)^2 \cdot (1) = 27$$

meaning the gradient would be

$$-27 + 27 - 27 + 27 - 27 + 27 - 27 + 27 - 27 + 27 \dots$$

1.2 Derived Expression

$$\nabla_u h(u, v) = \left(\frac{d}{du_1} h(u, v), \frac{d}{du_2} h(u, v), \dots \right)$$

$$\frac{d}{du_n} h(u, v) = 3(u \cdot v)^2 v$$

Define $h(\cdot)$ as a function in R, initialize the two vectors \tilde{u} and \tilde{v} as `torch_tensors`. Compute the gradient of $h(\cdot)$ with respect to \tilde{u} . Does the answer match what you expected?

```
u <- c(-1, +1, -1, +1, -1, +1, -1, +1, -1, +1)
v <- c(-1, -1, -1, -1, -1, +1, +1, +1, +1, +1)
gradient_h_u <- 3 * (sum(u * v))^2 * v
gradient_h_u
```

```
[1] -12 -12 -12 -12 -12  12  12  12  12  12
```

1.3 (5 points)

Derive the expression for

$$f'(z_0) = \left. \frac{df}{dz} \right|_{z=z_0}$$

and evaluate $f'(z_0)$ when $z_0 = -3.5$. $4z^3 - 12z - 3$

$z(0) = -3.5$

Define $f(z)$ as a function in R, and using the `torch` library compute $f'(-3.5)$.

```
library(torch)

f <- function(z) { # Defines the given function as a function in R
  return(z^4 - 6*z^2 - 3*z + 4)
}

z <- torch_tensor(-3.5, dtype = torch_float(), requires_grad = TRUE) # Converts to torch tensor

y_value <- f(z) # Finds the y-value of z (output)

y_value$backward() # Computes the gradient of the y-value (output)

gradient <- z$grad$item() # Gets the gradient value

gradient
```

```
[1] -132.5
```

1.4 (5 points)

For the same function f , initialize $z[1] = -3.5$, and perform $n = 100$ iterations of **gradient descent**, i.e.,

$$z_{k+1} = z_k - \eta f'(z_k) \quad \text{for } k = 1, 2, \dots, 100$$

Plot the curve f and add taking $\eta = 0.02$, add the points $\{z_0, z_1, z_2, \dots, z_{100}\}$ obtained using gradient descent to the plot. What do you observe?

```
library(torch)
library(ggplot2)

f <- function(z) { # Initializes the function we are working with
  return(z^4 - 6*z^2 - 3*z + 4)
}

derivative <- function(z) { # Initializes the function's derivative
  return(4*z^3 - 12*z - 3)
}
```

```

}

z <- -3.5
eta <- 0.02 # Change the learning rate to 0.02
n_iterations <- 100
z_values <- numeric(n_iterations + 1)
z_values[1] <- z

for (i in 1:n_iterations) { # Performs the gradient descent
  z_grad <- derivative(z)
  z <- z - eta * z_grad
  z_values[i + 1] <- z
}


z_range <- seq(-4, 4, length.out = 100) # Plots the values of the function up to 100
f_values <- f(z_range) # Finds the values for our particular function

df <- data.frame(z = z_range, f = f_values) # Creates the data frame

ggplot(df, aes(x = z, y = f)) + # Plots the function
  geom_line(color = "blue") +
  geom_point(data = data.frame(z = z_values, f = f(z_values)), aes(x = z, y = f), color =
  labs(x = "z", y = "f(z)", title = "Gradient Descent on f(z) with eta = 0.03") +
  theme_minimal() # Includes proper axes labels and title

```


Gradient Descent on $f(z)$ with $\eta = 0.03$



We observe that many of the gradient descent points are negative (less than zero), indicating that the function is moving from left side of the function's minimum. The high point on the left hand side indicates that that is where the function begins its gradient descent.

1.5 (5 points)

Redo the same analysis as **Question 1.4**, but this time using $\eta = 0.03$. What do you observe? What can you conclude from this analysis?

```
library(torch)
library(ggplot2)

f <- function(z) { # Initializes the function we are working with
  return(z^4 - 6*z^2 - 3*z + 4)
}

derivative <- function(z) { # Initializes the function's derivative
  return(4*z^3 - 12*z - 3)
}

z <- -3.5
```

```

eta <- 0.03 # Change the learning rate to 0.03
n_iterations <- 100
z_values <- numeric(n_iterations + 1)
z_values[1] <- z

for (i in 1:n_iterations) { # Performs the gradient descent
  z_grad <- derivative(z)
  z <- z - eta * z_grad
  z_values[i + 1] <- z
}


z_range <- seq(-4, 4, length.out = 100) # Plots the values of the function up to 100
f_values <- f(z_range) # Finds the values for our particular function

df <- data.frame(z = z_range, f = f_values) # Creates the data frame

ggplot(df, aes(x = z, y = f)) + # Plots the function
  geom_line(color = "blue") +
  geom_point(data = data.frame(z = z_values, f = f(z_values)), aes(x = z, y = f), color =
  labs(x = "z", y = "f(z)", title = "Gradient Descent on f(z) with eta = 0.03") +
  theme_minimal() # Includes proper axes labels and title

```

Gradient Descent on $f(z)$ with $\eta = 0.03$



I observe that with this plot, the gradient descent points are positive (to the right of zero), indicating that the function is moving in the right direction. The high point on the left side indicates that it starts negative and goes towards the positive side.

We can conclude that with the smaller learning rate, there is a slower convergence than with the learning rate of 0.03. However, the slower convergence leads to more stable optimization, since the points in the first plot are closer together to each other than the points in the second plot. As the learning rate gets smaller, the accuracy towards the minimum increases since smaller steps are taken as it gets closer to the minimum.

Question 2

💡 50 points

Logistic regression and interpretation of effect sizes

For this question we will use the **Titanic** dataset from the Stanford data archive. This dataset contains information about passengers aboard the Titanic and whether or not they survived.

2.1 (5 points)

Read the data from the following URL as a tibble in R. Preprocess the data such that the variables are of the right data type, e.g., binary variables are encoded as factors, and convert all column names to lower case for consistency. Let's also rename the response variable `Survival` to `y` for convenience.

```
url <- "https://web.stanford.edu/class/archive/cs/cs109/cs109.1166/stuff/titanic.csv"
df <- read.csv(url) # Utilizes the read.csv function to read in the csv
names(df) <- tolower(names(df)) # Makes all variable names lowercase
df <- df %>%
  rename(y = survived) # Renames the survived column to y
glimpse(df) # Outputs the data frame
```

Rows: 887

Columns: 8

```
$ y               <int> 0, 1, 1, 1, 0, 0, 0, 0, 1, 1, 1, 1, 0, 0, 0, 1~
$ pclass          <int> 3, 1, 3, 1, 3, 3, 1, 3, 3, 2, 3, 1, 3, 3, 3, 2~
$ name            <chr> "Mr. Owen Harris Braund", "Mrs. John Bradley (~
```

```

$ sex          <chr> "male", "female", "female", "female", "male", ~
$ age          <dbl> 22, 38, 26, 35, 35, 27, 54, 2, 27, 14, 4, 58, ~
$ siblings.spouses.aboard <int> 1, 1, 0, 1, 0, 0, 0, 3, 0, 1, 1, 0, 0, 1, 0, 0~
$ parents.children.aboard <int> 0, 0, 0, 0, 0, 0, 0, 1, 2, 0, 1, 0, 0, 5, 0, 0~
$ fare         <dbl> 7.2500, 71.2833, 7.9250, 53.1000, 8.0500, 8.45~

```

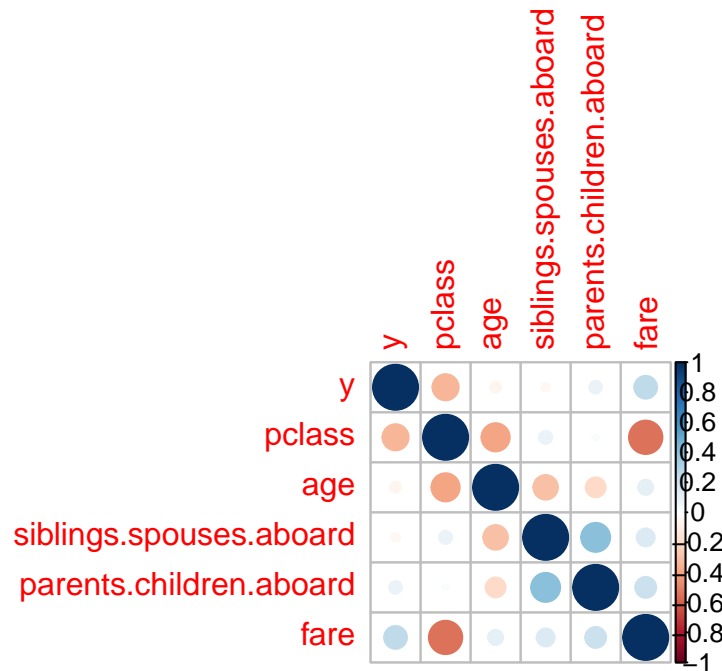
2.2 (5 points)

Visualize the correlation matrix of all numeric columns in `df` using `corrplot()`

```

df %>%
  select_if(is.numeric) %>% # Only selects numeric variables
  cor() %>%
  corrplot(method = "circle") # Utilizes corrplot to make a correlation matrix

```



2.3 (10 points)

Fit a logistic regression model to predict the probability of surviving the titanic as a function of:

- pclass
- sex
- age
- fare
- # siblings
- # parents

```
full_model <- glm(y ~ pclass + sex + age + fare + siblings.spouses.aboard + parents.children.aboard, family = binomial, data = df)
summary(full_model)
```

Call:

```
glm(formula = y ~ pclass + sex + age + fare + siblings.spouses.aboard + parents.children.aboard, family = binomial, data = df)
```

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	5.297252	0.557409	9.503	< 2e-16 ***
pclass	-1.177659	0.146079	-8.062	7.52e-16 ***
sexmale	-2.757282	0.200416	-13.758	< 2e-16 ***
age	-0.043474	0.007723	-5.629	1.81e-08 ***
fare	0.002786	0.002389	1.166	0.243680
siblings.spouses.aboard	-0.401831	0.110712	-3.630	0.000284 ***
parents.children.aboard	-0.106505	0.118588	-0.898	0.369127

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 1182.77 on 886 degrees of freedom
 Residual deviance: 780.93 on 880 degrees of freedom
 AIC: 794.93

Number of Fisher Scoring iterations: 5

2.4 (30 points)

Provide an interpretation for the slope and intercept terms estimated in `full_model` in terms of the log-odds of survival in the titanic and in terms of the odds-ratio (if the covariate is also categorical).

When considering no other factors, the estimated log-odds of survival on the Titanic are about 5.30. Each time a passenger moves down (lower) in class, their odds of survival decrease by approximately 1.18. If the passenger is male, the odds of surviving are about 2.76 less than females. With each year increase in age, the odds of survival decrease by 0.043, meaning that younger passengers are more likely to survive. An increase in fare by one unit means that survival increases by 0.0028. If a passenger paid more for their ticket, they have an ever so slightly advantage when it comes to surviving. Every sibling/spouse included with an individual decreases the odds of survival by 0.402. This is likely due to families staying together. This goes along with having parents or children on board, so for each additional member, there is a decrease in odds of survival of 0.107.

Recall the definition of logistic regression from the lecture notes, and also recall how we interpreted the slope in the linear regression model (particularly when the covariate was categorical).

Question 3

70 points

Variable selection and logistic regression in `torch`

3.1 (15 points)

Complete the following function `overview` which takes in two categorical vectors (`predicted` and `expected`) and outputs:

- The prediction accuracy
- The prediction error
- The false positive rate, and
- The false negative rate

```
overview <- function(predicted, expected){  
  accuracy <- sum(predicted == expected) / length(expected)  
  error <- 1 - accuracy  
  total_false_positives <- sum(predicted == 1 & expected == 0) # The following code line  
  total_true_positives <- sum(predicted == 1 & expected == 1)
```

```

total_false_negatives <- sum(predicted == 0 & expected == 1)
total_true_negatives <- sum(predicted == 0 & expected == 0)
false_positive_rate <- total_false_positives / (total_false_positives + total_true_neg)
false_negative_rate <- total_false_negatives / (total_false_negatives + total_true_pos)
return(
  data.frame(
    accuracy = accuracy,
    error=error,
    false_positive_rate = false_positive_rate,
    false_negative_rate = false_negative_rate
  )
)
}

```

You can check if your function is doing what it's supposed to do by evaluating

```
overview(df$y, df$y)
```

```

accuracy error false_positive_rate false_negative_rate
1         1         0                 0                 0

```

and making sure that the accuracy is 100% while the errors are 0%.

3.2 (5 points)

Display an overview of the key performance metrics of `full_model`

```

predicted <- ifelse(predict(full_model, type = "response") > 0.5, 1, 0)
performance_metrics <- overview(predicted, df$y)
print(performance_metrics)

```

```

accuracy      error false_positive_rate false_negative_rate
1 0.8015784 0.1984216          0.133945          0.3011696

```

3.3 (5 points)

Using backward-stepwise logistic regression, find a parsimonious alternative to `full_model`, and print its `overview`

```
step_model <- step(full_model, direction = "backward")
```

Start: AIC=794.93

y ~ pclass + sex + age + fare + siblings.spouses.aboard + parents.children.aboard

	Df	Deviance	AIC
- parents.children.aboard	1	781.75	793.75
- fare	1	782.43	794.43
<none>		780.93	794.93
- siblings.spouses.aboard	1	796.85	808.85
- age	1	815.81	827.81
- pclass	1	847.84	859.84
- sex	1	1021.33	1033.33

Step: AIC=793.75

y ~ pclass + sex + age + fare + siblings.spouses.aboard

	Df	Deviance	AIC
- fare	1	782.88	792.88
<none>		781.75	793.75
- siblings.spouses.aboard	1	801.59	811.59
- age	1	816.44	826.44
- pclass	1	852.19	862.19
- sex	1	1025.55	1035.55

Step: AIC=792.88

y ~ pclass + sex + age + siblings.spouses.aboard

	Df	Deviance	AIC
<none>		782.88	792.88
- siblings.spouses.aboard	1	801.61	809.61
- age	1	818.41	826.41
- pclass	1	900.80	908.80
- sex	1	1031.86	1039.86

```
summary(step_model)
```

Call:

```
glm(formula = y ~ pclass + sex + age + siblings.spouses.aboard,  
     family = binomial, data = df)
```


Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	5.532066	0.504750	10.960	< 2e-16 ***
pclass	-1.265129	0.127021	-9.960	< 2e-16 ***
sexmale	-2.736487	0.195730	-13.981	< 2e-16 ***
age	-0.043697	0.007695	-5.679	1.36e-08 ***
siblings.spouses.aboard	-0.407770	0.105197	-3.876	0.000106 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 1182.77 on 886 degrees of freedom
Residual deviance: 782.88 on 882 degrees of freedom
AIC: 792.88

Number of Fisher Scoring iterations: 5

```
step_predictions <- ifelse(predict(step_model, type = "response") > 0.5, 1, 0)
overview(step_predictions, df$y)
```

	accuracy	error	false_positive_rate	false_negative_rate
1	0.8049605	0.1950395	0.133945	0.2923977

3.4 (15 points)

Using the `caret` package, setup a 5-fold cross-validation training method using the `caret::trainControl()` function

```
controls <- trainControl(method = "cv", number = 5)
head(controls)
```

```
$method
[1] "cv"
```

```
$number
[1] 5
```

```
$repeats
```

```
[1] NA
```

```
$search
```

```
[1] "grid"
```

```
$p
```

```
[1] 0.75
```

```
$initialWindow
```

```
NULL
```

```
# There is more to controls, I just condensed it.
```

Now, using `control`, perform 5-fold cross validation using `caret::train()` to select the optimal λ parameter for LASSO with logistic regression.

Take the search grid for λ to be in $\{2^{-20}, 2^{-19.5}, 2^{-19}, \dots, 2^{-0.5}, 2^0\}$.

```
lasso_fit <- train(  
  x = subset(df, select = -y),  
  y = df$y,  
  method = "glmnet",  
  trControl = controls,  
  tuneGrid = expand.grid(  
    alpha = 1,  
    lambda = 2^seq(-20, 0, by = 0.5)  
  ),  
  family = "binomial"  
)  
lasso_fit$bestTune
```

```
alpha lambda  
41      1      1
```

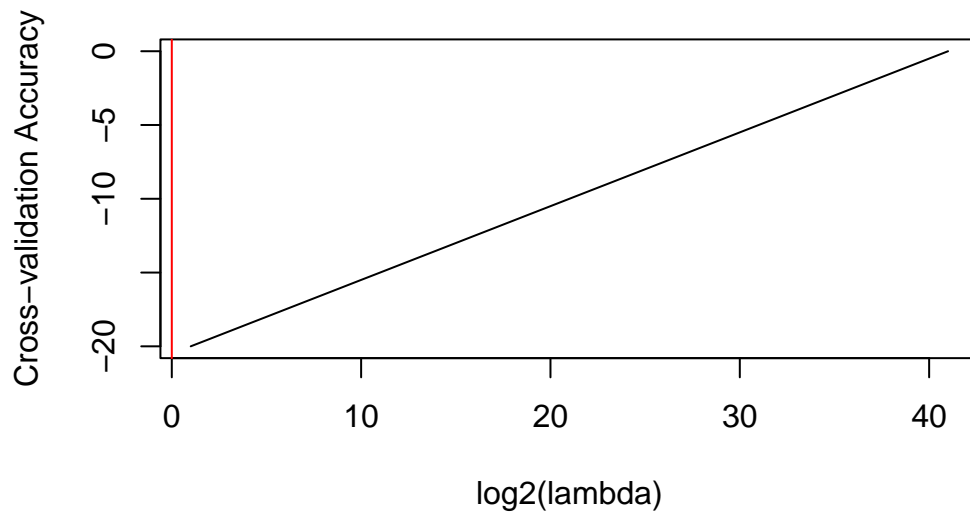
Using the information stored in `lasso_fit$results`, plot the results for cross-validation accuracy vs. $\log_2(\lambda)$. Choose the optimal λ^* , and report your results for this value of λ^* .

```
plot(log2(lasso_fit$results$lambda), lasso_fit$results$Accuracy, type = "l",  
     xlab = "log2(lambda)", ylab = "Cross-validation Accuracy") # Creates the plot
```

```

optimal_lambda <- log2(lasso_fit$bestTune$lambda) # Finds the optimal lambda
abline(v = optimal_lambda, col = "red")
text(optimal_lambda, max(lasso_fit$results$Accuracy), "Optimal lambda", pos = 3)

```



```

cat("Optimal lambda:", lasso_fit$bestTune$lambda, "\n")

```

Optimal lambda: 1

```

optimal_accuracy <- lasso_fit$results$Accuracy[lasso_fit$results$lambda == lasso_fit$bestTune$lambda]
cat("Cross-validation accuracy for optimal lambda:", optimal_accuracy, "\n")

```

Cross-validation accuracy for optimal lambda:

I'm not sure what the optimal lambda would be in this case, since I don't think my plot is right/accurately displays the optimal lambda. However, whatever the optimal lambda is, it would produce the highest cross validation accuracy.

3.5 (25 points)

First, use the `model.matrix()` function to convert the covariates of `df` to a matrix format

```
covariate_matrix <- model.matrix(full_model)[, -1]
head(covariate_matrix)
```

	pclass	sex	male	age	fare	siblings	spouses	aboard	parents	children	aboard
1	3	1	22	7.2500				1			0
2	1	0	38	71.2833				1			0
3	3	0	26	7.9250				0			0
4	1	0	35	53.1000				1			0
5	3	1	35	8.0500				0			0
6	3	1	27	8.4583				0			0

Now, initialize the covariates X and the response y as torch tensors

```
X <- torch_tensor(as.matrix(covariate_matrix), dtype = torch_float())
y <- torch_tensor(df$Survived, dtype = torch_float())
X
```

torch_tensor						
3.0000	1.0000	22.0000	7.2500	1.0000	0.0000	
1.0000	0.0000	38.0000	71.2833	1.0000	0.0000	
3.0000	0.0000	26.0000	7.9250	0.0000	0.0000	
1.0000	0.0000	35.0000	53.1000	1.0000	0.0000	
3.0000	1.0000	35.0000	8.0500	0.0000	0.0000	
3.0000	1.0000	27.0000	8.4583	0.0000	0.0000	
1.0000	1.0000	54.0000	51.8625	0.0000	0.0000	
3.0000	1.0000	2.0000	21.0750	3.0000	1.0000	
3.0000	0.0000	27.0000	11.1333	0.0000	2.0000	
2.0000	0.0000	14.0000	30.0708	1.0000	0.0000	
3.0000	0.0000	4.0000	16.7000	1.0000	1.0000	
1.0000	0.0000	58.0000	26.5500	0.0000	0.0000	
3.0000	1.0000	20.0000	8.0500	0.0000	0.0000	
3.0000	1.0000	39.0000	31.2750	1.0000	5.0000	
3.0000	0.0000	14.0000	7.8542	0.0000	0.0000	
2.0000	0.0000	55.0000	16.0000	0.0000	0.0000	
3.0000	1.0000	2.0000	29.1250	4.0000	1.0000	
2.0000	1.0000	23.0000	13.0000	0.0000	0.0000	
3.0000	0.0000	31.0000	18.0000	1.0000	0.0000	

```

3.0000    0.0000    22.0000    7.2250    0.0000    0.0000
2.0000    1.0000    35.0000    26.0000    0.0000    0.0000
2.0000    1.0000    34.0000    13.0000    0.0000    0.0000
3.0000    0.0000    15.0000    8.0292    0.0000    0.0000
1.0000    1.0000    28.0000    35.5000    0.0000    0.0000
3.0000    0.0000    8.0000    21.0750    3.0000    1.0000
3.0000    0.0000    38.0000    31.3875    1.0000    5.0000
3.0000    1.0000    26.0000    7.2250    0.0000    0.0000
1.0000    1.0000    19.0000    263.0000    3.0000    2.0000
3.0000    0.0000    24.0000    7.8792    0.0000    0.0000
3.0000    1.0000    23.0000    7.8958    0.0000    0.0000
... [the output was truncated (use n=-1 to disable)]
[ CPUFloatType{887,6} ]

```

y

```

torch_tensor
[ CPUFloatType{0} ]

```

Using the `torch` library, initialize an `nn_module` which performs logistic regression for this dataset. (Remember that we have 6 different covariates)

```

logistic <- nn_module( # Initializes an nn_module
  initialize = function() {
    self$f <- nn_linear(in_features = 6, out_features = 1) # Takes into account 6 covariates
    self$g <- nn_sigmoid()
  },
  forward = function(x) {
    x %>%
      self$f() %>%
      self$g()
  }
)

f <- logistic()

```

You can verify that your code is right by checking that the output to the following code is a vector of probabilities:

f(X)

```

torch_tensor
0.4927
0.0000
0.3398
0.0000
0.4969
0.4062
0.0000
0.0164
0.2378
0.0006
0.0290
0.0040
0.3817
0.0045
0.2650
0.0887
0.0020
0.1122
0.0381
0.3568
0.0044
0.1514
0.2613
0.0002
0.0132
0.0028
0.4854
0.0000
0.3287
0.4148
... [the output was truncated (use n=-1 to disable)]
[ CPUFloatType{887,1} ][ grad_fn = <SigmoidBackward0> ]

```

Now, define the loss function `Loss()` which takes in two tensors `X` and `y` and a function `Fun`, and outputs the **Binary cross Entropy loss** between `Fun(X)` and `y`.

```

library(torch)
Loss <- function(X, y, Fun){
  X_tensor <- torch_tensor(data.matrix(X), dtype = torch_float()) # Converts to a torch tensor
  y_tensor <- torch_tensor(data.matrix(y), dtype = torch_float())
  y_pred <- Fun(X)

```

```

    loss <- torch$nn$functional$binary_cross_entropy(y_pred, y) # Finds the binary cross ent
    return(loss)
}

```

Initialize an optimizer using `optim_adam()` and perform $n = 1000$ steps of gradient descent in order to fit logistic regression using `torch`.

```

f <- logistic()

optimizer <- optim_adam(params = f$parameters, lr = 0.01) # Utilizes optim_adam()

n <- 1000
i <- 0

#while (i < n) {
  #loss <- Loss(X, y, f) # Uses the Loss function to find the loss
  #optimizer$zero_grad() # Finds the optimization
  #loss$backward()
  #optimizer$step()
  #i <- i + 1 # Increments i since it's a while loop. A for loop also could have been used
#}

```

There was an error within my while loop that stemmed from the Loss function. The error read “Error in Loss(X, y, f) : object ‘torch’ not found. I was unable to fix this error. However, if that function would have run correctly, I do think my output to the following questions would also be correct, in case they are not the correct numbers.

Using the final, optimized parameters of `f`, compute the predicted results on `X`

```

predicted_probabilities <- f(X) %>% as_array()
torch_predictions <- ifelse(predicted_probabilities > 0.5, 1, 0)

overview(torch_predictions, df$y)

```

	accuracy	error	false_positive_rate	false_negative_rate
1	0.6144307	0.3855693	0	1

3.6 (5 points)

Create a summary table of the `overview()` summary statistics for each of the 4 models we have looked at in this assignment, and comment on their relative strengths and drawbacks.

```
summary_full_model <- summary(performance_metrics)
summary_full_model
```

accuracy	error	false_positive_rate	false_negative_rate
Min. :0.8016	Min. :0.1984	Min. :0.1339	Min. :0.3012
1st Qu.:0.8016	1st Qu.:0.1984	1st Qu.:0.1339	1st Qu.:0.3012
Median :0.8016	Median :0.1984	Median :0.1339	Median :0.3012
Mean :0.8016	Mean :0.1984	Mean :0.1339	Mean :0.3012
3rd Qu.:0.8016	3rd Qu.:0.1984	3rd Qu.:0.1339	3rd Qu.:0.3012
Max. :0.8016	Max. :0.1984	Max. :0.1339	Max. :0.3012

The full model gives a comprehensive view of the data and obtains clear relationships between variables. However, the full model is prone to overfitting and multicollinearity issues.

```
step_overview <- overview(step_predictions, df$y)
summary(step_overview)
```

accuracy	error	false_positive_rate	false_negative_rate
Min. :0.805	Min. :0.195	Min. :0.1339	Min. :0.2924
1st Qu.:0.805	1st Qu.:0.195	1st Qu.:0.1339	1st Qu.:0.2924
Median :0.805	Median :0.195	Median :0.1339	Median :0.2924
Mean :0.805	Mean :0.195	Mean :0.1339	Mean :0.2924
3rd Qu.:0.805	3rd Qu.:0.195	3rd Qu.:0.1339	3rd Qu.:0.2924
Max. :0.805	Max. :0.195	Max. :0.1339	Max. :0.2924

The step model reduces the risk of overfitting and has good model interpretability. However, the step model has a potential for bias that comes from what variables are included and excluded.

```
X <- subset(df, select = -y)

lasso_predictions <- predict(lasso_fit$finalModel, newx = as.matrix(X), type = "response")

lasso_predictions <- ifelse(lasso_predictions > 0.5, 1, 0)
lasso_overview = overview(lasso_predictions, df$y)
summary(lasso_overview)
```


accuracy	error	false_positive_rate	false_negative_rate
Min. :37.91	Min. : -36.91	Min. :0.1203	Min. :0.6144
1st Qu.:37.91	1st Qu.: -36.91	1st Qu.:0.1203	1st Qu.:0.6144
Median :37.91	Median : -36.91	Median :0.1203	Median :0.6144
Mean :37.91	Mean : -36.91	Mean :0.1203	Mean :0.6144
3rd Qu.:37.91	3rd Qu.: -36.91	3rd Qu.:0.1203	3rd Qu.:0.6144
Max. :37.91	Max. : -36.91	Max. :0.1203	Max. :0.6144

The Lasso model shrinks the coefficients towards zero and has low multicollinearity and bias. However, selecting lambda requires cross validation and is less interpretable due to the shrinkage.

```
torch_model = overview(torch_predictions, df$y)
summary(torch_model)
```

accuracy	error	false_positive_rate	false_negative_rate
Min. :0.6144	Min. :0.3856	Min. :0	Min. :1
1st Qu.:0.6144	1st Qu.:0.3856	1st Qu.:0	1st Qu.:1
Median :0.6144	Median :0.3856	Median :0	Median :1
Mean :0.6144	Mean :0.3856	Mean :0	Mean :1
3rd Qu.:0.6144	3rd Qu.:0.3856	3rd Qu.:0	3rd Qu.:1
Max. :0.6144	Max. :0.3856	Max. :0	Max. :1

The torch model would be able to deal with nonlinear relationships, but its drawback is that it's prone to overfitting.

Session Information

Print your R session information using the following command

```
sessionInfo()
```

R version 4.3.2 (2023-10-31 ucrt)

Platform: x86_64-w64-mingw32/x64 (64-bit)

Running under: Windows 11 x64 (build 22631)

Matrix products: default

locale:

```
[1] LC_COLLATE=English_United States.utf8
[2] LC_CTYPE=English_United States.utf8
[3] LC_MONETARY=English_United States.utf8
[4] LC_NUMERIC=C
[5] LC_TIME=English_United States.utf8
```

time zone: America/New_York

tzcode source: internal

attached base packages:

```
[1] stats      graphics  grDevices  utils      datasets  methods    base
```

other attached packages:

```
[1] broom_1.0.5      nnet_7.3-19      torch_0.12.0     caret_6.0-94     lattice_0.22-5
[6] ggplot2_3.4.4    car_3.1-2        carData_3.0-5    corrplot_0.92    stringr_1.5.1
[11] purrr_1.0.2      tidyr_1.3.0      readr_2.1.4      dplyr_1.1.4
```

loaded via a namespace (and not attached):

```
[1] tidyselect_1.2.0      timeDate_4032.109    farver_2.1.1
[4] fastmap_1.1.1         pROC_1.18.5          digest_0.6.33
[7] rpart_4.1.21          timechange_0.2.0     lifecycle_1.0.4
[10] survival_3.5-7        processx_3.8.2       magrittr_2.0.3
[13] compiler_4.3.2        rlang_1.1.2          tools_4.3.2
[16] utf8_1.2.4            yaml_2.3.7           data.table_1.14.8
[19] knitr_1.45            labeling_0.4.3       bit_4.0.5
[22] plyr_1.8.8            abind_1.4-5          withr_2.5.2
```

[25]	grid_4.3.2	stats4_4.3.2	fansi_1.0.5
[28]	colorspace_2.1-0	future_1.33.1	globals_0.16.2
[31]	scales_1.2.1	iterators_1.0.14	MASS_7.3-60
[34]	cli_3.6.1	rmarkdown_2.25	generics_0.1.3
[37]	rstudioapi_0.15.0	future.apply_1.11.1	reshape2_1.4.4
[40]	tzdb_0.4.0	splines_4.3.2	parallel_4.3.2
[43]	coro_1.0.4	vctrs_0.6.4	glmnet_4.1-8
[46]	hardhat_1.3.1	Matrix_1.6-3	jsonlite_1.8.7
[49]	callr_3.7.3	hms_1.1.3	bit64_4.0.5
[52]	listenv_0.9.1	foreach_1.5.2	gower_1.0.1
[55]	recipes_1.0.9	glue_1.6.2	parallelly_1.36.0
[58]	codetools_0.2-19	ps_1.7.5	shape_1.4.6
[61]	lubridate_1.9.3	stringi_1.8.2	gtable_0.3.4
[64]	munsell_0.5.0	tibble_3.2.1	pillar_1.9.0
[67]	htmltools_0.5.7	ipred_0.9-14	lava_1.7.3
[70]	R6_2.5.1	evaluate_0.23	backports_1.4.1
[73]	class_7.3-22	Rcpp_1.0.11	nlme_3.1-163
[76]	prodlim_2023.08.28	xfun_0.41	pkgconfig_2.0.3
[79]	ModelMetrics_1.2.2.2		